

NAME:- DUBEY KARAN SANJEEV
CLASS:- B.E - 4
ROLL NO:- 04
BATCH:- A

EXPERIMENT NO-1

AIM: a: To study overview of Hadoop Ecosystem tools.

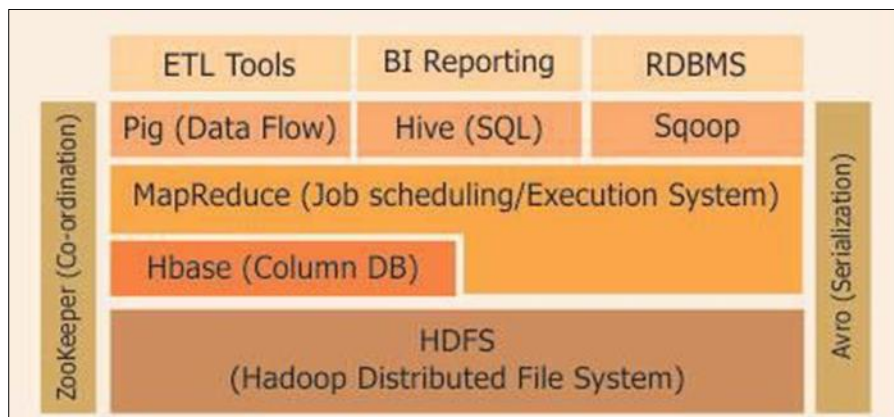
b: To execute basic commands of Hadoop eco system components like HDFS, Hive, Sqoop.

THEORY:

What is Hadoop?

Hadoop is an open-source software framework for storing data and running applications on clusters of commodity hardware. It provides massive storage for any kind of data, enormous processing power and the ability to handle virtually limitless concurrent tasks or jobs.

Hadoop Ecosystem



Hadoop is distributed master slave architecture that consist of

1. Hadoop distributed file system (HDFS) for storage.

- It is storage component of Hadoop.
- It is distributed file system that is modeled after Google File System (GFS).
- It is optimized for high throughput and works best while reading and writing large files.

2. Map Reduce for computational capabilities.

- It is the heart of Hadoop.
- It provides a mechanism for programmer to control distributed systems for processing data sets.

1. Pig:

- A data flow language and execution environment for exploring very large datasets.
- Pig is made up of two pieces:
 - The language used to express data flows, called Pig Latin.
 - The execution environment to run Pig Latin programs.

- There are currently two environments: local execution in a single JVM and distributed execution on a Hadoop cluster.
- A Pig Latin program is made up of a series of operations, or transformations, that are applied to the input data to produce output.
- Pig is a scripting language for exploring large datasets, Writing the mappers and reducers, compiling and packaging the code, submitting the job(s), and retrieving the results is a time consuming task.
- Pig has the ability to process terabytes of data simply by issuing half dozen times of pig latin from console.

2.Hive:

- A distributed data warehouse.
- Hive manages data stored in HDFS and provides a query language based on SQL (and which is translated by the runtime engine to MapReduce jobs) for querying the data.
- Hive was created to make it possible for analysts with strong SQL skills (but meager Java programming skills) to run queries on the huge volumes of data

3.Sqoop

- A tool for efficient bulk transfer of data between structured data stores such as RDBMS and HDFS.
- Hadoop platform has ability to work with data in several forms.
- Often, valuable data in an organization is stored in structured data stores such as RDBMS.
- Sqoop is an open-source tool that allows users to extract data from a relational database into Hadoop for further processing.

4. HBase

- A distributed, column-oriented database. HBase uses HDFS for its underlying storage, and supports both batch-style computations using MapReduce and point queries (random reads).
- HBase is not relational and does not support SQL, but given the proper problem space, it is able to do what an RDBMS cannot such as host very large, sparsely populated tables on clusters made from commodity hardware.

5.Zookeeper:

- A distributed, highly available coordination service. ZooKeeper provides primitives such as distributed locks that can be used for building distributed applications.
- Writing distributed applications is hard.
- For example

When a message is sent across the network between two nodes and the network fails, the sender does not know whether the receiver got the message. It may have gotten through before the network failed, or it may not have. Or perhaps the receiver's process died. The only way that the sender can find out what happened is to reconnect to the receiver and ask it. This is partial failure: when we don't even know if an operation failed. ZooKeeper can't make partial failures go away, since they are intrinsic to distributed systems. It certainly does not hide partial failures, either.¹ But what ZooKeeper does do is give you a set of tools to build distributed applications that can safely handle partial failures.

6.Avro

- A serialization system for efficient, cross-language RPC, and persistent data storages.
- Apache Avro is a language-neutral data serialization system.
- Avro provides
 - Rich data structures.

- A compact, fast, binary data format.
- A container file, to store a persistent data.
- RPC.

CONCLUSION:

Hadoop ecosystem consists of variety technologies which can be used for various applications and basic commands of HDFS are executed which are used for map reduce execution.

Program formation/ Execution/ ethical practices (06)	Timely Submission and Documentation (02)	Viva Answer (02)	Experiment Marks (10)	Teacher Signature with date

1) help HDFS Shell Command

Syntax of help hdfs Command

```
$ hadoop fs -help
```

Help hdfs shell command helps hadoop developers figure out all the available hadoop commands and how to use them.



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -help  
Usage: hadoop fs [generic options]  
    [-cat [-ignoreCrc] <src> ...]  
    [-chgrp [-R] GROUP PATH...]  
    [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]  
    [-chown [-R] [OWNER][:[GROUP]] PATH...]  
    [-copyFromLocal <localsrc> ... <dst>]  
    [-copyToLocal [-ignoreCrc] [-crc] <src> ... <localdst>]  
    [-count [-q] <path> ...]  
    [-cp <src> ... <dst>]  
    [-df [-h] [<path> ...]]  
    [-du [-s] [-h] <path> ...]  
    [-expunge]  
    [-get [-ignoreCrc] [-crc] <src> ... <localdst>]  
    [-getmerge [-nl] <src> <localdst>]  
    [-help [cmd ...]]  
    [-ls [-d] [-h] [-R] [<path> ...]]  
    [-mkdir [-p] <path> ...]  
    [-moveFromLocal <localsrc> ... <dst>]  
    [-moveToLocal <src> <localdst>]  
    [-mv <src> ... <dst>]  
    [-put <localsrc> ... <dst>]  
    [-rm [-f] [-r|-R] [-skipTrash] <src> ...]  
    [-rmdir [--ignore-fail-on-non-empty] <dir> ...]  
    [-setrep [-R] [-w] <rep> <path/file> ...]  
    [-stat [format] <path> ...]  
    [-tail [-f] <file>]  
    [-test [-ezd] <path>]  
    [-text [-ignoreCrc] <src> ...]  
    [-touchz <path> ...]  
    [-usage [cmd ...]]  
  
-cat [-ignoreCrc] <src> ...:    Fetch all files that match the file pattern <src>  
>  
                                and display their content on stdout.  
  
-chgrp [-R] GROUP PATH...:      This is equivalent to -chown ... :GROUP ...
```

Variations of the Hadoop fs Help Command

```
$ hadoop fs -help ls
```

Using the help command with a specific command lists the usage information along with the options to use the command.

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -help ls  
-ls [-d] [-h] [-R] [<path> ...]: List the contents that match the specifi  
ed file pattern. If  
    path is not specified, the contents of /user/<currentUser>  
    will be listed. Directory entries are of the form  
        dirName (full path) <dir>  
    and file entries are of the form  
        fileName(full path) <r n> size  
    where n is the number of replicas specified for the file  
    and size is the size of the file, in bytes.  
    -d Directories are listed as plain files.  
    -h Formats the sizes of files in a human-readable fashion  
        rather than a number of bytes.  
    -R Recursively list the contents of directories.  
[cloudera@localhost ~]$
```

2) Usage HDFS Shell Command

```
$ hadoop fs -usage ls
```

Usage command gives all the options that can be used with a particular hdfs command.

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -usage ls  
Usage: hadoop fs [generic options] -ls [-d] [-h] [-R] [<path> ...]  
[cloudera@localhost ~]$
```

3) ls HDFS Shell Command

Syntax for ls Hadoop Command -

```
$ hadoop fs -ls
```

This command will list all the available files and subdirectories under default directory. For instance, in our example the default directory for Cloudera VM is /user/cloudera

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls  
Found 4 items  
drwx----- - cloudera cloudera      0 2016-09-06 19:00 .staging  
drwxr-xr-x - cloudera cloudera      0 2016-09-06 18:50 Input  
drwxr-xr-x - cloudera cloudera      0 2016-09-06 19:00 Output  
drwxr-xr-x - cloudera cloudera      0 2016-08-25 19:53 tweets  
[cloudera@localhost ~]$
```

Variations of Hadoop ls Shell Command

```
$ hadoop fs -ls /
```

Returns all the available files and subdirectories present under the root directory.

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls /  
Found 5 items  
drwxr-xr-x - hbase hbase 0 2013-10-07 08:19 /hbase  
drwxr-xr-x - solr solr 0 2013-10-07 08:18 /solr  
drwxrwxrwx - hdfs supergroup 0 2013-10-07 08:18 /tmp  
drwxr-xr-x - hdfs supergroup 0 2013-10-07 08:20 /user  
drwxr-xr-x - hdfs supergroup 0 2013-10-07 08:18 /var  
[cloudera@localhost ~]$
```

\$ hadoop fs -ls -R /user/cloudera

Returns all the available files and recursively lists all the subdirectories under /user/Cloudera

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls -R /user/cloudera  
drwx----- - cloudera cloudera 0 2016-09-06 19:00 /user/cloudera/.sta  
ging  
drwxr-xr-x - cloudera cloudera 0 2016-09-06 18:50 /user/cloudera/Inpu  
t  
-rw-r--r-- 3 cloudera cloudera 3291648 2016-09-06 18:50 /user/cloudera/Inpu  
t/war_and_peace  
drwxr-xr-x - cloudera cloudera 0 2016-09-06 19:00 /user/cloudera/Outp  
ut  
-rw-r--r-- 3 cloudera cloudera 0 2016-09-06 19:00 /user/cloudera/Outp  
ut/SUCCESS  
drwxr-xr-x - cloudera cloudera 0 2016-09-06 18:57 /user/cloudera/Outp  
ut/logs  
drwxr-xr-x - cloudera cloudera 0 2016-09-06 19:00 /user/cloudera/Outp  
ut/logs/history  
-rw-r--r-- 3 cloudera cloudera 23558 2016-09-06 19:00 /user/cloudera/Outp  
ut/logs/history/job_201609061101_0001_1473213466855_cloudera_WordCount  
-rw-r--r-- 3 cloudera cloudera 75212 2016-09-06 18:57 /user/cloudera/Outp  
ut/logs/history/job_201609061101_0001_conf.xml  
-rw-r--r-- 3 cloudera cloudera 467842 2016-09-06 19:00 /user/cloudera/Outp  
ut/part-00000  
drwxr-xr-x - cloudera cloudera 0 2016-08-25 19:53 /user/cloudera/twee  
ts  
-rw-r--r-- 3 cloudera cloudera 2464 2016-08-25 19:53 /user/cloudera/twee  
ts/flume-twitter-partioned.conf  
-rw-r--r-- 3 cloudera cloudera 1369 2016-08-25 19:53 /user/cloudera/twee  
ts/flume.conf  
-rw-r--r-- 3 cloudera cloudera 1487 2016-08-25 19:53 /user/cloudera/twee  
ts/twitter-part.conf  
[cloudera@localhost ~]$
```

4) mkdir- Used to create a new directory in HDFS at a given location.

Example of HDFS mkdir Command -

\$ hadoop fs -mkdir /user/cloudera/dezyre1

The above command will create a new directory named dezyre1 under the location /user/cloudera



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -mkdir /user/cloudera/dezyre1  
[cloudera@localhost ~]$
```

Note : Cloudera and other [hadoop distribution vendors](#) provide /user/ directory with read/write permission to all users but other directories are available as read-only. Thus, to create a folder in the root directory, users require superuser permission as shown below -

\$ sudo -u hdfs hadoop fs -mkdir /dezyre

This command will create a new directory named dezyre under the / (root directory).

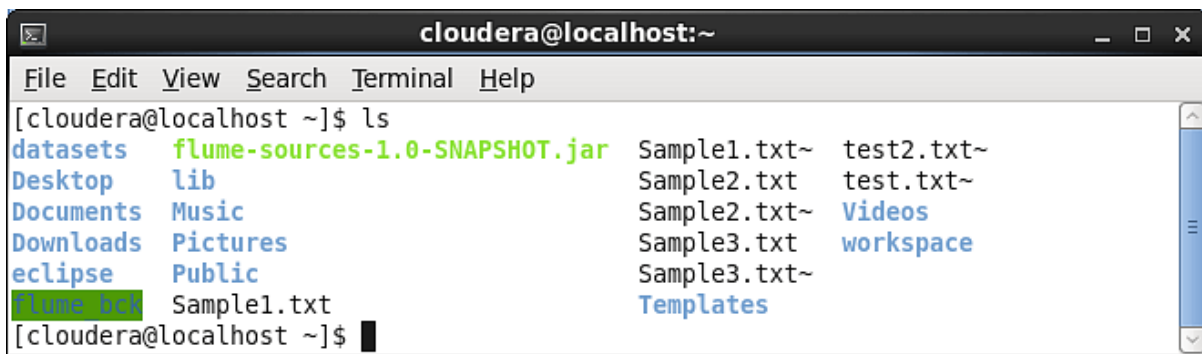


```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ sudo -u hdfs hadoop fs -mkdir /dezyre  
[cloudera@localhost ~]$
```

5) copyFromLocal

Copy a file from local filesystem to HDFS location.

For the following examples, we will use Sample.txt file available in the /home/Cloudera location.



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ ls  
datasets  flume-sources-1.0-SNAPSHOT.jar  Sample1.txt~  test2.txt~  
Desktop   lib                               Sample2.txt~  test.txt~  
Documents Music                           Sample2.txt~  Videos  
Downloads Pictures                       Sample3.txt   workspace  
eclipse   Public                           Sample3.txt~  
flume.txt Sample1.txt                       Templates  
[cloudera@localhost ~]$
```

Example - \$ hadoop fs -copyFromLocal Sample1.txt /user/cloudera/dezyre1

Copy/Upload Sample1.txt available in /home/cloudera (local default) to /user/cloudera/dezyre1 (hdfs path)




```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -copyFromLocal Sample1.txt /user/cloudera/dezyre1  
[cloudera@localhost ~]$
```

6) put -

This hadoop command uploads a single file or multiple source files from local file system to hadoop distributed file system (HDFS).

Ex - \$ hadoop fs -put Sample2.txt /user/cloudera/dezyre1

Copy/Upload Sample2.txt available in /home/cloudera (local default) to /user/cloudera/dezyre1 (hdfs path)



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -put Sample2.txt /user/cloudera/dezyre1  
[cloudera@localhost ~]$
```

7) moveFromLocal

This hadoop command functions similar to the put command but the source file will be deleted after copying.

Example - \$ hadoop fs -moveFromLocal Sample3.txt /user/cloudera/dezyre1

Move Sample3.txt available in /home/cloudera (local default) to /user/cloudera/dezyre1 (hdfs path). Source file will be deleted after moving.



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -moveFromLocal Sample3.txt /user/cloudera/dezyre1  
[cloudera@localhost ~]$
```



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ ls  
datasets  eclipse  Music  Sample1.txt~  Videos  
Desktop    flume-1.0.0  Pictures  Sample2.txt  workspace  
Documents  flume-sources-1.0-SNAPSHOT.jar  Public  Sample2.txt~  
Downloads  lib        Sample1.txt  Templates  
[cloudera@localhost ~]$
```

8) du

Displays the disk usage for all the files available under a given directory.

Example - \$ hadoop fs -du /user/cloudera/dezyre1


```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -du /user/cloudera/dezyre1  
370 /user/cloudera/dezyre1/Sample1.txt  
370 /user/cloudera/dezyre1/Sample2.txt  
370 /user/cloudera/dezyre1/Sample3.txt  
[cloudera@localhost ~]$
```

9) df

Display disk usage of current hadoop distributed file system.

Example - `$ hadoop fs -df`

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -df  
Filesystem                                Size      Used    Available  Use%  
hdfs://localhost.localdomain:8020 28799380685 91496448 24875388928    0%  
[cloudera@localhost ~]$
```

10) Expunge

This HDFS command empties the trash by deleting all the files and directories.

Example - `$ hadoop fs -expunge`

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -expunge  
[cloudera@localhost ~]$
```

11) Cat

This is similar to the cat command in Unix and displays the contents of a file.

Example - `$ hadoop fs -cat /user/cloudera/dezyre1/Sample1.txt`

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -cat /user/cloudera/dezyre1/Sample1.txt  
-----  
/ This is Sample-1 file. /  
-----  
|| Created to explain the functionality of Hadoop FsShell ||  
-----  
+++++  
| Thanks for choosing dezyre. |  
+++++  
[cloudera@localhost ~]$
```

12) cp

Copy files from one HDFS location to another HDFS location.

Example - \$ hadoop fs -cp /user/cloudera/dezyre/war_and_peace /user/cloudera/dezyre1/

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -cp /user/cloudera/dezyre/war_and_peace /user/  
cloudera/dezyre1/  
[cloudera@localhost ~]$
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/dezyre  
Found 1 items  
-rw-r--r-- 3 cloudera cloudera 3291648 2016-09-29 14:19 /user/cloudera/dezy  
re/war_and_peace  
[cloudera@localhost ~]$
```

13) mv

Move files from one HDFS location to another HDFS location.

Example - \$ hadoop fs -mv /user/cloudera/dezyre1/Sample1.txt /user/cloudera/dezyre/

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/dezyrel  
Found 4 items  
-rw-r--r--  3 cloudera cloudera      370 2016-09-29 13:14 /user/cloudera/dezyrel/Sample1.txt  
-rw-r--r--  3 cloudera cloudera      370 2016-09-29 13:16 /user/cloudera/dezyrel/Sample2.txt  
-rw-r--r--  3 cloudera cloudera      370 2016-09-29 13:17 /user/cloudera/dezyrel/Sample3.txt  
-rw-r--r--  3 cloudera cloudera 3291648 2016-09-29 14:22 /user/cloudera/dezyrel/war_and_peace  
[cloudera@localhost ~]$
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/dezyre  
Found 1 items  
-rw-r--r--  3 cloudera cloudera 3291648 2016-09-29 14:19 /user/cloudera/dezyre/war_and_peace  
[cloudera@localhost ~]$
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -mv /user/cloudera/dezyrel/Sample1.txt /user/cloudera/dezyre/  
[cloudera@localhost ~]$
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/dezyre  
Found 2 items  
-rw-r--r--  3 cloudera cloudera      370 2016-09-29 13:14 /user/cloudera/dezyre/Sample1.txt  
-rw-r--r--  3 cloudera cloudera 3291648 2016-09-29 14:19 /user/cloudera/dezyre/war_and_peace  
[cloudera@localhost ~]$
```

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/dezyrel  
Found 3 items  
-rw-r--r--  3 cloudera cloudera      370 2016-09-29 13:16 /user/cloudera/dezyrel/Sample2.txt  
-rw-r--r--  3 cloudera cloudera      370 2016-09-29 13:17 /user/cloudera/dezyrel/Sample3.txt  
-rw-r--r--  3 cloudera cloudera 3291648 2016-09-29 14:22 /user/cloudera/dezyrel/war_and_peace  
[cloudera@localhost ~]$
```

14) rm

Removes the file or directory from the mentioned HDFS location.

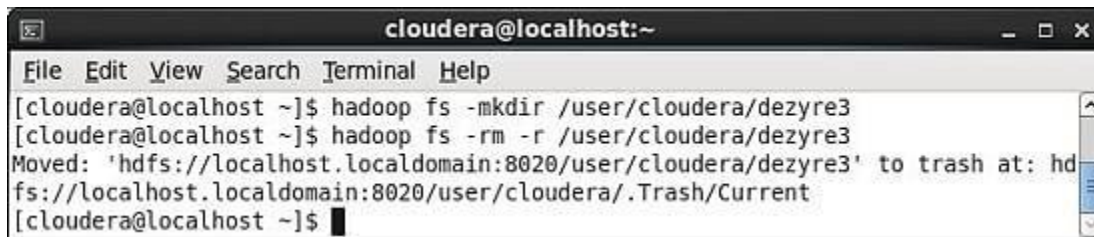
Example – `$ hadoop fs -rm -r /user/cloudera/dezyre3`



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -mkdir /user/cloudera/dezyre3  
[cloudera@localhost ~]$ hadoop fs -rm -r /user/cloudera/dezyre3  
Moved: 'hdfs://localhost.localdomain:8020/user/cloudera/dezyre3' to trash at: hdfs://localhost.localdomain:8020/user/cloudera/.Trash/Current  
[cloudera@localhost ~]$
```

Example – `$ hadoop fs -rm -r /user/cloudera/dezyre3`


Deletes or removes the directory and its content from HDFS location in a recursive manner.



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -mkdir /user/cloudera/dezyre3  
[cloudera@localhost ~]$ hadoop fs -rm -r /user/cloudera/dezyre3  
Moved: 'hdfs://localhost.localdomain:8020/user/cloudera/dezyre3' to trash at: hdfs://localhost.localdomain:8020/user/cloudera/.Trash/Current  
[cloudera@localhost ~]$
```

Example – `$ hadoop fs -rm /user/cloudera/dezyre3`

Delete or remove the files from HDFS location.

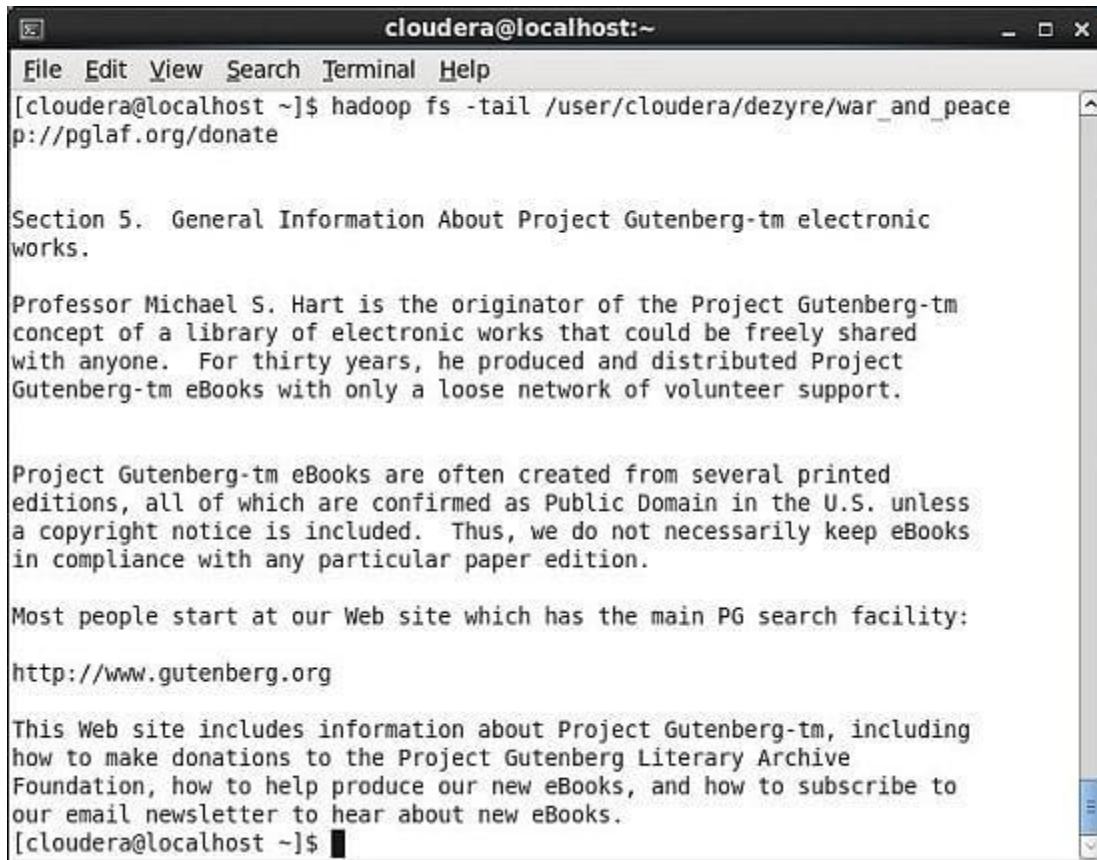


```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -rm /user/cloudera/dezyre1/war_and_peace  
Moved: 'hdfs://localhost.localdomain:8020/user/cloudera/dezyre1/war_and_peace' to trash at: hdfs://localhost.localdomain:8020/user/cloudera/.Trash/Current  
[cloudera@localhost ~]$
```

15) tail

This hadoop command will show the last kilobyte of the file to stdout.

Example – `$ hadoop fs -tail /user/cloudera/dezyre/war_and_peace`

A terminal window titled 'cloudera@localhost:~' with a menu bar (File, Edit, View, Search, Terminal, Help). The command '[cloudera@localhost ~]\$ hadoop fs -tail /user/cloudera/dezyre/war_and_peace' has been executed. The output shows the last kilobyte of the file in a page-wise format, including sections on Project Gutenberg-tm electronic works, the originator Michael S. Hart, and the Project Gutenberg website.

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -tail /user/cloudera/dezyre/war_and_peace  
p://pglaf.org/donate  
  
Section 5. General Information About Project Gutenberg-tm electronic  
works.  
  
Professor Michael S. Hart is the originator of the Project Gutenberg-tm  
concept of a library of electronic works that could be freely shared  
with anyone. For thirty years, he produced and distributed Project  
Gutenberg-tm eBooks with only a loose network of volunteer support.  
  
Project Gutenberg-tm eBooks are often created from several printed  
editions, all of which are confirmed as Public Domain in the U.S. unless  
a copyright notice is included. Thus, we do not necessarily keep eBooks  
in compliance with any particular paper edition.  
  
Most people start at our Web site which has the main PG search facility:  
  
http://www.gutenberg.org  
  
This Web site includes information about Project Gutenberg-tm, including  
how to make donations to the Project Gutenberg Literary Archive  
Foundation, how to help produce our new eBooks, and how to subscribe to  
our email newsletter to hear about new eBooks.  
[cloudera@localhost ~]$
```

Example – \$ hadoop fs -tail -f /user/cloudera/dezyre/war_and_peace

Using the tail commands with -f option, shows the last kilobyte of the file from end in a page wise format.


```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -tail -f /user/cloudera/dezyre/war_and_peace  
p://pglaf.org/donate  
  
Section 5. General Information About Project Gutenberg-tm electronic  
works.  
  
Professor Michael S. Hart is the originator of the Project Gutenberg-tm  
concept of a library of electronic works that could be freely shared  
with anyone. For thirty years, he produced and distributed Project  
Gutenberg-tm eBooks with only a loose network of volunteer support.  
  
Project Gutenberg-tm eBooks are often created from several printed  
editions, all of which are confirmed as Public Domain in the U.S. unless  
a copyright notice is included. Thus, we do not necessarily keep eBooks  
in compliance with any particular paper edition.  
  
Most people start at our Web site which has the main PG search facility:  
  
http://www.gutenberg.org  
  
This Web site includes information about Project Gutenberg-tm, including  
how to make donations to the Project Gutenberg Literary Archive  
Foundation, how to help produce our new eBooks, and how to subscribe to  
our email newsletter to hear about new eBooks.  
█
```

16) copyToLocal

Copies the files to the local filesystem . This is similar to `hadoop fs -get` command but in this case the destination location must be a local file reference

Example - `$ hadoop fs -copyFromLocal /user/cloudera/dezyre1/Sample1.txt /home/cloudera/hdfs_bkp/`

Copy/Download Sample1.txt available in /user/cloudera/dezyre1 (hdfs path) to /home/cloudera/hdfs_bkp/ (local path)

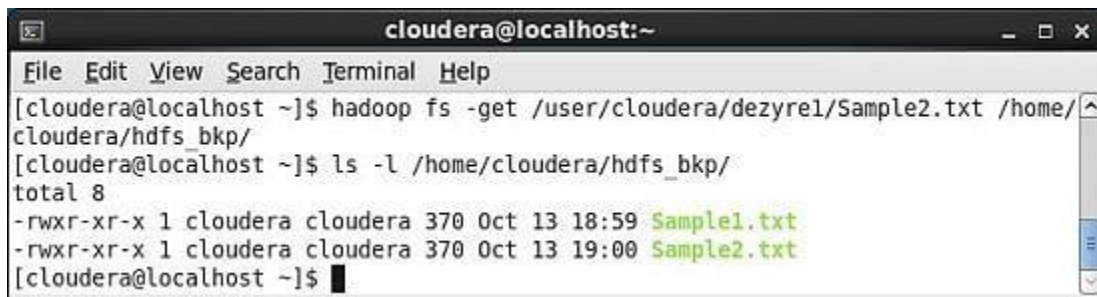
```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -copyToLocal /user/cloudera/dezyre1/Sample1.tx  
t /home/cloudera/hdfs_bkp/  
[cloudera@localhost ~]$ ls -l /home/cloudera/hdfs_bkp/  
total 4  
-rwxr-xr-x 1 cloudera cloudera 370 Oct 13 18:59 Sample1.txt  
[cloudera@localhost ~]$ █
```

17) get

Downloads or Copies the files to the local filesystem.

Example - `$ hadoop fs -get /user/cloudera/dezyre1/Sample2.txt /home/cloudera/hdfs_bkp/`

Copy/Download Sample2.txt available in /user/cloudera/dezyre1 (hdfs path) to /home/cloudera/hdfs_bkp/ (local path)



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -get /user/cloudera/dezyre1/Sample2.txt /home/  
cloudera/hdfs_bkp/  
[cloudera@localhost ~]$ ls -l /home/cloudera/hdfs_bkp/  
total 8  
-rwxr-xr-x 1 cloudera cloudera 370 Oct 13 18:59 Sample1.txt  
-rwxr-xr-x 1 cloudera cloudera 370 Oct 13 19:00 Sample2.txt  
[cloudera@localhost ~]$
```

18) touchz

Used to create an empty file at the specified location.

Example - `$ hadoop fs -touchz /user/cloudera/dezyre1/Sample4.txt`

It will create a new empty file Sample4.txt in /user/cloudera/dezyre1/ (hdfs path)



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -touchz /user/cloudera/dezyre1/Sample4.txt  
[cloudera@localhost ~]$ hadoop fs -ls /user/cloudera/dezyre1/  
Found 4 items  
-rw-r--r-- 3 cloudera cloudera 370 2016-09-29 17:48 /user/cloudera/dezy  
re1/Sample1.txt  
-rw-r--r-- 3 cloudera cloudera 370 2016-09-29 13:16 /user/cloudera/dezy  
re1/Sample2.txt  
-rw-r--r-- 3 cloudera cloudera 370 2016-09-29 13:17 /user/cloudera/dezy  
re1/Sample3.txt  
-rw-r--r-- 3 cloudera cloudera 0 2016-10-13 19:10 /user/cloudera/dezy  
re1/Sample4.txt  
[cloudera@localhost ~]$
```

19) setrep

This hadoop fs command is used to set the replication for a specific file.

Example - `$ hadoop fs -setrep -w 1 /user/cloudera/dezyre1/Sample1.txt`

It will set the replication factor of Sample1.txt to 1



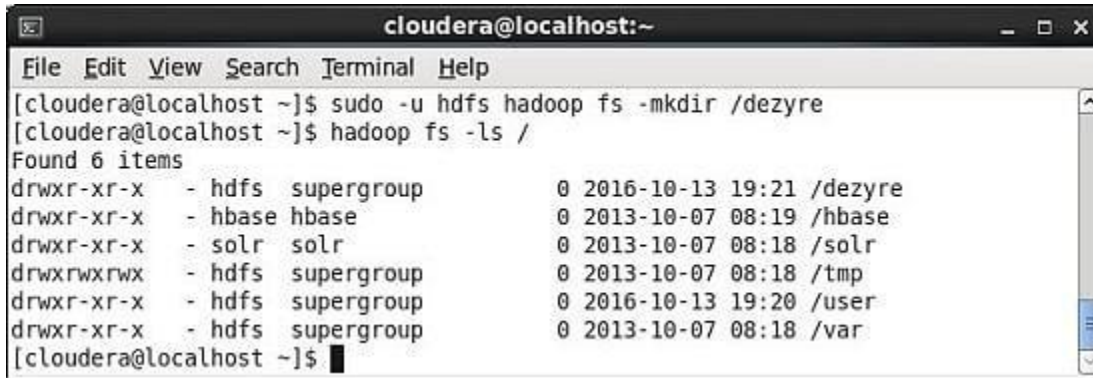
```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -setrep -w 1 /user/cloudera/dezyre1/Sample1.tx  
t  
Replication 1 set: /user/cloudera/dezyre1/Sample1.txt  
Waiting for /user/cloudera/dezyre1/Sample1.txt ... done  
[cloudera@localhost ~]$
```

20) chgrp

This hadoop command is basically used to change the group name.

Example - `$ sudo -u hdfs hadoop fs -chgrp -R cloudera /dezyre`

It will change the /dezyre directory group membership from supergroup to cloudera (To perform this operation superuser permission is required)



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ sudo -u hdfs hadoop fs -mkdir /dezyre  
[cloudera@localhost ~]$ hadoop fs -ls /  
Found 6 items  
drwxr-xr-x - hdfs supergroup          0 2016-10-13 19:21 /dezyre  
drwxr-xr-x - hbase hbase              0 2013-10-07 08:19 /hbase  
drwxr-xr-x - solr solr                0 2013-10-07 08:18 /solr  
drwxrwxrwx - hdfs supergroup          0 2013-10-07 08:18 /tmp  
drwxr-xr-x - hdfs supergroup          0 2016-10-13 19:20 /user  
drwxr-xr-x - hdfs supergroup          0 2013-10-07 08:18 /var  
[cloudera@localhost ~]$
```



```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ sudo -u hdfs hadoop fs -chgrp -R cloudera /dezyre  
[cloudera@localhost ~]$ hadoop fs -ls /dezyre  
[cloudera@localhost ~]$ hadoop fs -ls /  
Found 6 items  
drwxr-xr-x - hdfs cloudera            0 2016-10-13 19:21 /dezyre  
drwxr-xr-x - hbase hbase              0 2013-10-07 08:19 /hbase  
drwxr-xr-x - solr solr                0 2013-10-07 08:18 /solr  
drwxrwxrwx - hdfs supergroup          0 2013-10-07 08:18 /tmp  
drwxr-xr-x - hdfs supergroup          0 2016-10-13 19:20 /user  
drwxr-xr-x - hdfs supergroup          0 2013-10-07 08:18 /var  
[cloudera@localhost ~]$
```

21) chown

This command lets you change both the owner and group name simultaneously.

Example - `$ sudo -u hdfs hadoop fs -chown -R cloudera /dezyre`

It will change the /dezyre directory ownership from hdfs user to cloudera user (To perform this operation superuser is permission required)

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ sudo -u hdfs hadoop fs -chown -R cloudera /dezyre  
[cloudera@localhost ~]$ hadoop fs -ls /  
Found 6 items  
drwxr-xr-x - cloudera cloudera 0 2016-10-13 19:21 /dezyre  
drwxr-xr-x - hbase hbase 0 2013-10-07 08:19 /hbase  
drwxr-xr-x - solr solr 0 2013-10-07 08:18 /solr  
drwxrwxrwx - hdfs supergroup 0 2013-10-07 08:18 /tmp  
drwxr-xr-x - hdfs supergroup 0 2016-10-13 19:20 /user  
drwxr-xr-x - hdfs supergroup 0 2013-10-07 08:18 /var  
[cloudera@localhost ~]$
```

22) hadoop chmod

Used to change the permissions of a given file/dir.

Example - \$ hadoop fs -chmod /dezyre

It will change the /dezyre directory permission to 700 (drwx-----).

```
cloudera@localhost:~  
File Edit View Search Terminal Help  
[cloudera@localhost ~]$ hadoop fs -chmod 700 /dezyre  
[cloudera@localhost ~]$ hadoop fs -ls /  
Found 6 items  
drwx----- - cloudera cloudera 0 2016-10-13 19:21 /dezyre  
drwxr-xr-x - hbase hbase 0 2013-10-07 08:19 /hbase  
drwxr-xr-x - solr solr 0 2013-10-07 08:18 /solr  
drwxrwxrwx - hdfs supergroup 0 2013-10-07 08:18 /tmp  
drwxr-xr-x - hdfs supergroup 0 2016-10-13 19:20 /user  
drwxr-xr-x - hdfs supergroup 0 2013-10-07 08:18 /var  
[cloudera@localhost ~]$
```

Note : hadoop chmod 777

To execute this , the user must be the owner of the file or must be a super user. On executing this command, all users will get read,write and execute permission on the file.